

Predictive study of the Titanic incident

November 1, 2023

1 Predictive analysis

Based on certain characteristics of the passengers of the Titanic, we seek to build a classification algorithm that can predict with excellent efficiency the survival case of a passenger of the Titanic.

To do this, we use the different Scikit-Learn pipelines to preprocess the data, find the best hyper-parameters and find the best classification algorithm among those tested.

[Data Source](#)

1.0.1 Importing the basics libraries

```
[163]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
pd.set_option("display.notebook_repr_html", False)
```

1.0.2 Importing dataset

```
[178]: df = pd.read_csv("Titanic-Dataset.csv")

print("Data shape :", df.shape)

print("\nTwo first row of the dataset :\n")

df.head()
```

Data shape : (891, 12)

Two first row of the dataset :

```
[178]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

		Name	Sex	Age	SibSp	\
0		Braund, Mr. Owen Harris	male	22.0	1	
1	Cummings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	1	
2		Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)		female	35.0	1	
4		Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

1.0.3 Data info

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age            714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch          891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Cabin           204 non-null   object
11  Embarked        889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

1.1 Data preprocessing

SibSp:

This feature is based on the assumption that for a passenger whose loved ones have survived, the chances of survival are higher, and vice versa. Given the database provided, it cannot be said that there is a relationship between two passengers. We therefore prefer to remove this variable, but if you want more information, we advise you to take a look at this publication via this link below:

<https://www.kaggle.com/code/ailuropus/extracting-family-relationships-on-titanic-sibsp>

Name :

We prefer to remove this variable to simplify our work, but if you want to use it, you can use sklearn's CountVectorizer module to convert this variable to numbers. This will help determine if the name had an impact on a deceased person's case.

Ticket :

We also remove this variable, it deserves special treatment and the base dataset does not provide much information about it.

Passenger ID:

This feature is not important for predictive modeling. We delete it too.

1.1.1 Missing values

```
[4]: df.isnull().sum()
```

```
[4]: PassengerId      0
     Survived        0
     Pclass          0
     Name            0
     Sex             0
     Age            177
     SibSp           0
     Parch           0
     Ticket          0
     Fare            0
     Cabin          687
     Embarked        2
     dtype: int64
```

1.1.2 Let's put the features in a list according to their context to facilitate their pretreatment

```
[5]: cat_missing_values = ['Embarked', 'Cabin']

     cat_without_missing = ['Sex']

     num_missing_values = ['Age']

     drop_columns = ['PassengerId', 'Ticket', 'Name']

     num_columns = ['Pclass', 'SibSp', 'Parch', 'Fare']
```

1.1.3 Importing librairies

```
[6]: from sklearn.compose import make_column_transformer
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.pipeline import make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split
```

1.1.4 Pipeline to replace missing categorical values and convert categories into numbers

```
[7]: cat_preprocessor = make_pipeline(
    ↳ SimpleImputer(strategy='constant', missing_values=np.nan, ↳
    ↳ fill_value='missing'),
    OneHotEncoder(handle_unknown='ignore', ↳
    ↳ sparse_output=False)
)
```

1.1.5 Pipeline for scaling numeric values

```
[8]: num_preprocessor = make_pipeline(SimpleImputer(strategy="median"), ↳
    ↳ MinMaxScaler())
```

1.1.6 Final data preprocessing

```
[9]: data_preprocessing = make_column_transformer(
    (OneHotEncoder(sparse_output=False), cat_without_missing),
    (cat_preprocessor, cat_missing_values),
    (num_preprocessor, num_missing_values),
    ('drop', drop_columns),
    (MinMaxScaler(), num_columns)
)
```

1.1.7 Split the dataset into x and y

```
[10]: X = df.drop("Survived", axis = 1)
y = df["Survived"]
```

```
[11]: X.shape
```

```
[11]: (891, 11)
```

```
[12]: y.shape
```

```
[12]: (891,)
```

1.1.8 Fit data_preprocessing

```
[14]: data_array = data_preprocessing.fit_transform(X)
```

```
[179]: df_encoded = pd.DataFrame(data_array, columns=data_preprocessing.  
    ↪ get_feature_names_out())  
df_encoded.head(2)
```

```
[179]: onehotencoder__Sex_female  onehotencoder__Sex_male  pipeline-1__Embarked_C  \  
0                                0.0                      1.0                      0.0  
1                                1.0                      0.0                      1.0  
  
    pipeline-1__Embarked_Q  pipeline-1__Embarked_S  \  
0                          0.0                      1.0  
1                          0.0                      0.0  
  
    pipeline-1__Embarked_missing  pipeline-1__Cabin_A10  pipeline-1__Cabin_A14  \  
0                                0.0                      0.0                      0.0  
1                                0.0                      0.0                      0.0  
  
    pipeline-1__Cabin_A16  pipeline-1__Cabin_A19  ...  pipeline-1__Cabin_F38  \  
0                          0.0                      0.0  ...                      0.0  
1                          0.0                      0.0  ...                      0.0  
  
    pipeline-1__Cabin_F4  pipeline-1__Cabin_G6  pipeline-1__Cabin_T  \  
0                          0.0                      0.0                      0.0  
1                          0.0                      0.0                      0.0  
  
    pipeline-1__Cabin_missing  pipeline-2__Age  minmaxscaler__Pclass  \  
0                              1.0           0.271174                1.0  
1                              0.0           0.472229                0.0  
  
    minmaxscaler__SibSp  minmaxscaler__Parch  minmaxscaler__Fare  
0                    0.125                    0.0           0.014151  
1                    0.125                    0.0           0.139136  
  
[2 rows x 159 columns]
```

1.1.9 Split dataset with train_test_split method

```
[16]: X_train, X_test, y_train, y_test = train_test_split(  
                                             X, y, test_size=0.25,  
                                             random_state=42)
```

```
[17]: X_train.shape
```

```
[17]: (668, 11)
```

```
[18]: X_test.shape
```

```
[18]: (223, 11)
```

1.2 Choose the classification algorithm using GridSearchCV method of Sklearn

Selection of the most efficient classifier with the technical GridSearchCV

1.2.1 Random Forest Classifier

```
[19]: from sklearn.ensemble import RandomForestClassifier as rfc
```

```
[47]: pipe_rfc = Pipeline(steps=[('preprocessor', data_preprocessing),  
                                ('rf_classifier', rfc(random_state = 42))])
```

```
[48]: param_dict = {  
        'rf_classifier__n_estimators' : [5, 10, 15, 20, 30, 60, 80, 100],  
        'rf_classifier__max_features': ['sqrt', 'log2', None, .1, .25, .3, .35, .4],  
        'rf_classifier__max_depth' : [None, 4, 7, 10, 15, 20, 25, 30, 35],  
        'rf_classifier__criterion': ['gini', 'entropy', 'log_loss']  
    }  
param_dict
```

```
[48]: {'rf_classifier__n_estimators': [5, 10, 15, 20, 30, 60, 80, 100],  
      'rf_classifier__max_features': ['sqrt',  
      'log2',  
      None,  
      0.1,  
      0.25,  
      0.3,  
      0.35,  
      0.4],  
      'rf_classifier__max_depth': [None, 4, 7, 10, 15, 20, 25, 30, 35],  
      'rf_classifier__criterion': ['gini', 'entropy', 'log_loss']}
```

1.2.2 Cross validation

```
[49]: from sklearn.model_selection import KFold

cross_validation= KFold(n_splits=5,
                        shuffle=True,
                        random_state=42)

cross_validation
```

```
[49]: KFold(n_splits=5, random_state=42, shuffle=True)
```

1.2.3 GridSearch

```
[50]: from sklearn.model_selection import GridSearchCV
```

```
[51]: GridSear_rfc = GridSearchCV(pipe_rfc, param_dict, cv = cross_validation)
```

```
[52]: GridSear_rfc
```

```
[52]: GridSearchCV(cv=KFold(n_splits=5, random_state=42, shuffle=True),
                  estimator=Pipeline(steps=[('preprocessor',
ColumnTransformer(transformers=[('onehotencoder',
OneHotEncoder(sparse_output=False),
['Sex']),
('pipeline-1',
Pipeline(steps=[('simpleimputer',
SimpleImputer(fill_value='missing',
strategy='constant')),
('onehotencoder',
OneHotEncoder(handle_unknown='i...
['Pclass',
'SibSp',
'Parch',
'Fare'])))]),
                  ('rf_classifier',
RandomForestClassifier(random_state=42)))]),
                  param_grid={'rf_classifier__criterion': ['gini', 'entropy',
'log_loss'],
'rf_classifier__max_depth': [None, 4, 7, 10, 15, 20,
25, 30, 35],
'rf_classifier__max_features': ['sqrt', 'log2', None,
0.1, 0.25, 0.3, 0.35,
0.4],
'rf_classifier__n_estimators': [5, 10, 15, 20, 30, 60,
80, 100]})
```

```
[53]: GridSear_rfc.fit(X_train, y_train)
```

```
[53]: GridSearchCV(cv=KFold(n_splits=5, random_state=42, shuffle=True),
                  estimator=Pipeline(steps=[('preprocessor',
ColumnTransformer(transformers=[('onehotencoder',
OneHotEncoder(sparse_output=False),
['Sex'])),
('pipeline-1',
Pipeline(steps=[('simpleimputer',
SimpleImputer(fill_value='missing',
strategy='constant')),
('onehotencoder',
OneHotEncoder(handle_unknown='i...
['Pclass',
'SibSp',
'Parch',
'Fare']]]))),
                  ('rf_classifier',
RandomForestClassifier(random_state=42))]),
                  param_grid={'rf_classifier__criterion': ['gini', 'entropy',
'log_loss'],
'rf_classifier__max_depth': [None, 4, 7, 10, 15, 20,
25, 30, 35],
'rf_classifier__max_features': ['sqrt', 'log2', None,
0.1, 0.25, 0.3, 0.35,
0.4],
'rf_classifier__n_estimators': [5, 10, 15, 20, 30, 60,
80, 100]})
```

```
[54]: GridSear_rfc.best_params_
```

```
[54]: {'rf_classifier__criterion': 'entropy',
'rf_classifier__max_depth': 10,
'rf_classifier__max_features': 0.35,
'rf_classifier__n_estimators': 15}
```

```
[196]: round(GridSear_rfc.best_score_*100, 4)
```

```
[196]: 84.279
```

1.2.4 K-Nearest Neighbors Classifier

```
[56]: from sklearn.neighbors import KNeighborsClassifier
```

```
[57]: pipe_knn = Pipeline(steps=[('preprocessor', data_preprocessing),
('KN_classifier', KNeighborsClassifier())])
```

```
[58]:
```



```

params_knn = {'KN_classifier__n_neighbors': [1, 3, 5, 7, 9, 11, 12, 13, 14, 15, 16, 18, 20] ,
              "KN_classifier__weights": ['uniform', 'distance'],
              "KN_classifier__algorithm": ['auto', 'ball_tree', 'kd_tree', 'brute'],
              "KN_classifier__leaf_size" : [10, 20, 30, 40, 50, 60, 70, 80, 90, 100]
              }

params_knn

```

```

[58]: {'KN_classifier__n_neighbors': [1, 3, 5, 7, 9, 11, 12, 13, 14, 15, 16, 18, 20],
      'KN_classifier__weights': ['uniform', 'distance'],
      'KN_classifier__algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
      'KN_classifier__leaf_size': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100]}

```

```

[59]: GridSear_knn = GridSearchCV(pipe_knn, params_knn , cv = cross_validation)

```

```

[60]: GridSear_knn.fit(X_train, y_train)

```

```

[60]: GridSearchCV(cv=KFold(n_splits=5, random_state=42, shuffle=True),
                  estimator=Pipeline(steps=[('preprocessor',
      ColumnTransformer(transformers=[('onehotencoder',
      OneHotEncoder(sparse_output=False),
      ['Sex'])),
      ('pipeline-1',
      Pipeline(steps=[('simpleimputer',
          SimpleImputer(fill_value='missing',
                        strategy='constant')),
          ('onehotencoder',
          OneHotEncoder(handle_unknown='i...
      ('minmaxscaler',
      MinMaxScaler(),
      ['Pclass',
      'SibSp',
      'Parch',
      'Fare']])))),
      ('KN_classifier',
      KNeighborsClassifier()))],
      param_grid={'KN_classifier__algorithm': ['auto', 'ball_tree',
      'kd_tree', 'brute'],
      'KN_classifier__leaf_size': [10, 20, 30, 40, 50, 60,
      70, 80, 90, 100],
      'KN_classifier__n_neighbors': [1, 3, 5, 7, 9, 11, 12,
      13, 14, 15, 16, 18,
      20],
      'KN_classifier__weights': ['uniform', 'distance']})

```

```
[61]: GridSear_knn.best_params_
```

```
[61]: {'KN_classifier__algorithm': 'auto',  
      'KN_classifier__leaf_size': 10,  
      'KN_classifier__n_neighbors': 18,  
      'KN_classifier__weights': 'uniform'}
```

```
[194]: round(GridSear_knn.best_score_*100, 4)
```

```
[194]: 80.6834
```

1.2.5 Adaboost classifier

```
[63]: # https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html
```

```
from sklearn.ensemble import AdaBoostClassifier
```

```
[64]: pipe_adb = Pipeline(steps=[('preprocessor', data_preprocessing),  
                                ('adb_classifier', AdaBoostClassifier(random_state =  
                                ↪42))])
```

```
[65]: params_adb = {'adb_classifier__n_estimators': [100, 120, 150, 160, 180, 200,  
            ↪250, 300, 350] ,  
                  "adb_classifier__learning_rate": [0.1, 0.2, 0.25, 0.3, 0.35, 0.  
            ↪4, 0.5, 0.7, 0.8, 0.9, 1],  
                  "adb_classifier__algorithm": ['SAMME', 'SAMME.R']  
                  }
```

```
params_adb
```

```
[65]: {'adb_classifier__n_estimators': [100, 120, 150, 160, 180, 200, 250, 300, 350],  
      'adb_classifier__learning_rate': [0.1,  
      0.2,  
      0.25,  
      0.3,  
      0.35,  
      0.4,  
      0.5,  
      0.7,  
      0.8,  
      0.9,  
      1],  
      'adb_classifier__algorithm': ['SAMME', 'SAMME.R']}
```

```
[66]: GridSear_adb = GridSearchCV(pipe_adb, params_adb , cv = cross_validation)
```

```
[67]: GridSear_adb.fit(X_train, y_train)
```

```
[67]: GridSearchCV(cv=KFold(n_splits=5, random_state=42, shuffle=True),
                  estimator=Pipeline(steps=[('preprocessor',
ColumnTransformer(transformers=[('onehotencoder',
OneHotEncoder(sparse_output=False),
['Sex']]),
('pipeline-1',
Pipeline(steps=[('simpleimputer',
SimpleImputer(fill_value='missing',
strategy='constant')),
('onehotencoder',
OneHotEncoder(handle_unknown='i...
'Ticket',
'Name']]),
('minmaxscaler',
MinMaxScaler(),
['Pclass',
'SibSp',
'Parch',
'Fare']]])),
('adb_classifier',
AdaBoostClassifier(random_state=42))]),
param_grid={'adb_classifier__algorithm': ['SAMME', 'SAMME.R'],
'adb_classifier__learning_rate': [0.1, 0.2, 0.25, 0.3,
0.35, 0.4, 0.5, 0.7,
0.8, 0.9, 1],
'adb_classifier__n_estimators': [100, 120, 150, 160,
180, 200, 250, 300,
350]})
```

```
[68]: GridSear_adb.best_params_
```

```
[68]: {'adb_classifier__algorithm': 'SAMME',
'adb_classifier__learning_rate': 1,
'adb_classifier__n_estimators': 300}
```

```
[195]: round(GridSear_adb.best_score_*100, 4)
```

```
[195]: 81.4252
```

1.2.6 Hist Gradient Boosting Classifier

```
[80]: from sklearn.ensemble import HistGradientBoostingClassifier as HGBClassifier
```

```
[81]: pipe_hgb = Pipeline(steps=[('preprocessor', data_preprocessing),
('hgb_classifier', HGBClassifier(random_state = 42))])
```

```
[105]: params_hgb = {'hgb_classifier__learning_rate': [0.03, 0.04, 0.05, 0.05, 0.1, 0.2, 0.25],
                    "hgb_classifier__l2_regularization" : [0, 4, 8, 10, 16, 20, 25],
                    "hgb_classifier__max_depth" : [6, 7, 8, 9, 12, 14, None],
                    }
```

```
params_hgb
```

```
[105]: {'hgb_classifier__learning_rate': [0.02, 0.03, 0.4, 0.05, 0.1, 0.2, 0.25],
        'hgb_classifier__l2_regularization': [0, 4, 8, 10, 16, 20, 25],
        'hgb_classifier__max_depth': [6, 7, 8, 9, 12, 14, None]}
```

```
[106]: GridSear_hgb = GridSearchCV(pipe_hgb, params_hgb , cv = cross_validation)
```

```
[107]: GridSear_hgb.fit(X_train, y_train)
```

```
[107]: GridSearchCV(cv=KFold(n_splits=5, random_state=42, shuffle=True),
                    estimator=Pipeline(steps=[('preprocessor',
ColumnTransformer(transformers=[('onehotencoder',
OneHotEncoder(sparse_output=False),
['Sex']),
('pipeline-1',
Pipeline(steps=[('simpleimputer',
SimpleImputer(fill_value='missing',
strategy='constant')),
('onehotencoder',
OneHotEncoder(handle_unknown='i...
['PassengerId',
'Ticket',
'Name'])),
('minmaxscaler',
MinMaxScaler(),
['Pclass',
'SibSp',
'Parch',
'Fare']]])),
                    ('hgb_classifier',
HistGradientBoostingClassifier(random_state=42))]),
                    param_grid={'hgb_classifier__l2_regularization': [0, 4, 8, 10, 16,
20, 25],
'hgb_classifier__learning_rate': [0.02, 0.03, 0.4,
0.05, 0.1, 0.2,
0.25],
'hgb_classifier__max_depth': [6, 7, 8, 9, 12, 14,
None]})
```

```
[108]: GridSear_hgb.best_params_
```

```
[197]: round(GridSear_hgb.best_score_*100, 4)
```

1.3 Selection of the best model

```
[129]:
```

	RF Classifier	KNN Classifier	ADB Classifier	\
GridSearch_best_score	84.28	80.68	81.43	
	HGB Classifier			
GridSearch best score	83.97			

The Random Forest Classifier achieved an efficiency of over 83.825% using the GridSearch method. We select this model and train it on the data set with the best hyperparameters

1.3.1 We train the best pipeline on all our data

```
[131]: final_model = Pipeline([('data_preprocessing', data_preprocessing),
                                ('rfc', rfc(random_state=42,
                                             max_depth=7,
                                             max_features='sqrt',
                                             n_estimators=100))
                                ])

final_model.fit(X, y)
```

[illegible]

```

('onehotencoder',
OneHotEncoder(handle_unknown='ignore',
sparse_output=False))),

['Embarked', 'Cabin']],

('pipeline-2',

Pipeline(steps=[('simpleimputer',
SimpleImputer(strategy='median')),
('minmaxscaler',
MinMaxScaler())]),

['Age']],

('drop', 'drop',
['PassengerId', 'Ticket',
'Name']),

('minmaxscaler',
MinMaxScaler(),
['Pclass', 'SibSp', 'Parch',
'Fare']))),

('rfc', RandomForestClassifier(max_depth=7, random_state=42)))]

```

```
[198]: round(final_model.score(X, y)*100, 4)
```

```
[198]: 86.0831
```

Our model is ready to be saved for future predictions or for an internet deployment.

1.4 Save model

```
[141]: import joblib
```

```
[142]: # save the model to a file
joblib.dump(final_model, 'randomforest_classifier.joblib')
```

```
[142]: ['randomforest_classifier.joblib']
```

1.5 Prediction

Predict a small random sample from the dataset.

```
[145]: random_set = df.sample(n=4, random_state=42)

random_set_without_target = random_set.drop("Survived", axis = 1)

random_set_without_target
```

```
[145]:
```

	PassengerId	Pclass	Name \
709	710	3	Moubarek, Master. Halim Gonios ("William George")
439	440	2	Kvillner, Mr. Johan Henrik Johannesson
840	841	3	Alhomaki, Mr. Ilmari Rudolf

720		721	2		Harper, Miss. Annie Jessie "Nina"
	Sex	Age	SibSp	Parch	
709	male	NaN	1	1	2661 15.2458 NaN C
439	male	31.0	0	0	C.A. 18723 10.5000 NaN S
840	male	20.0	0	0	SOTON/02 3101287 7.9250 NaN S
720	female	6.0	0	1	248727 33.0000 NaN S

```
[146]: model_save = joblib.load('randomforest_classifier.joblib')
print(model_save)
```

```
Pipeline(steps=[('data_preprocessing',
                  ColumnTransformer(transformers=[('onehotencoder',
                                                    OneHotEncoder(sparse_output=False),
                                                                    ['Sex']),
                                                    ('pipeline-1',
                                                     Pipeline(steps=[('simpleimputer',
                                                                       SimpleImputer(fill_value='missing',
                                                                       strategy='constant')),
                                                                    ('onehotencoder',
                                                                       OneHotEncoder(handle_unknown='ignore',
                                                                       sparse_output=False))])),
                                                                    ['Embarked', 'Cabin']),
                                                    ('pipeline-2',
                                                     Pipeline(steps=[('simpleimputer',
                                                                       SimpleImputer(strategy='median')),
                                                                    ('minmaxscaler',
                                                                       MinMaxScaler()))],
                                                                    ['Age']),
                                                                    ('drop', 'drop',
                                                                       ['PassengerId', 'Ticket',
                                                                       'Name']),
                                                                    ('minmaxscaler',
                                                                       MinMaxScaler(),
                                                                       ['Pclass', 'SibSp', 'Parch',
                                                                       'Fare'])])),
                  ('rfc', RandomForestClassifier(max_depth=7, random_state=42))])
```

```
[147]: model_save.predict(random_set_without_target)
```

```
[147]: array([0, 0, 0, 1], dtype=int64)
```

1.5.1 The true values that were predicted

```
[148]: np.array(random_set["Survived"])
```

```
[148]: array([1, 0, 0, 1], dtype=int64)
```